

Project NeuroTone: Comprehensive Project Report

A Non-Invasive, Smartphone-Compatible Voice Acoustic Screening Prototype for Parkinson's Disease

Document Version: 1.0 (Formal)	Current Status: Phase 3 Completed (Benchmark & Visual Deck)
Project Lead: NeuroTone Research & Engineering Team	Target Audience: Medical, Technical, and Executive Stakeholders

1. Executive Summary & Project Vision

Parkinson's Disease (PD) is among the most prevalent neurodegenerative conditions globally, characterized by the progressive loss of dopamine-generating neurons in the brain. Historically, clinical diagnosis depends on observable physical motor signs, notably resting tremors, limb rigidity, and postural instability. However, by the time these overt motor symptoms become clinically apparent, an estimated 60% to 80% of dopamine-producing cells in the substantia nigra have already been irreversibly lost. Early, accessible, and continuous screening is therefore critical to facilitate timely medical and therapeutic intervention.

Project NeuroTone addresses this diagnostic gap by turning a ubiquitous device—a standard smartphone—into an instantaneous, non-invasive early screening tool. Up to 90% of individuals living with Parkinson's experience subtle vocal degradation (clinically termed *hypophonia* and *laryngeal dysarthria*) in early disease stages. By prompting a patient to sustain a simple vowel phonation (such as saying 'ah' for three seconds into a phone microphone), NeuroTone extracts micro-acoustic biomarkers of vocal fold tension and neurological stability, generating an objective risk score in seconds without expensive specialized hospital hardware.

2. Clinical & Biological Foundations: The Voice as a Neurological Window

Human voice production requires sub-millisecond neurological synchronization across the brainstem, laryngeal vocal folds, respiratory diaphragm, and oral articulation tract. In Parkinson's Disease, basal ganglia dysfunction impairs fine motor control, producing three distinctive acoustic phenomena:

- **Laryngeal Muscle Rigidity & Incomplete Closure:** Stiffened vocal cord muscles cannot fully seal during vocal vibration. This allows turbulent unphonated air to escape, creating breathiness, roughness, and diminished tonal clarity.
- **Neurological Vocal Micro-Tremors:** Involuntary micro-oscillations in laryngeal muscle tension introduce rapid, cycle-to-cycle instability in fundamental frequency (pitch jitter) and wave volume (amplitude shimmer).
- **Chaotic Signal Turbulence:** Healthy vocal fold vibration is smooth and harmonic. Parkinsonian vocal production displays elevated physical entropy and nonlinear turbulence, which can be quantified mathematically.

Because vocal changes frequently precede observable hand tremors by months or years, telephonic and mobile acoustic screening provides an accessible, non-stigmatizing gateway for timely clinical evaluation.

3. System Architecture & Engineering Organization

NeuroTone is built according to professional Python engineering standards, ensuring end-to-end reproducibility, preventing training data leakage, and enabling modular plug-and-play model iteration.

Directory / File	Engineering Role & Contents
data/raw/	Immutable raw datasets: UCI Parkinson's benchmark, plus reserved folders for Italian and NeuroVoz corpora.
data/processed/	Standardized, cleaned datasets (uci_cleaned.csv) stripped of non-predictive string identifiers.
src/	Core pipeline modules: data downloaders, cleaning scripts, and modular cross-validation evaluation engines.
src/plotter/	Visualization engine and clinical feature mapping (feature_mapping.py) for readable reporting.
visualizations/	High-resolution pitch-deck visual assets (300 DPI PNGs) and compressed distribution archives.
docs/	Documentation for context retention (project_context.md) and formal reports (report.pdf).
requirements.txt	Pinned dependencies across scientific computing, audio processing, tree ensembles, and reporting.

4. Dataset Profile & Exploratory Analysis

For Phase 1 through Phase 3, we evaluated the benchmark **UCI Parkinson's Disease Dataset**, compiled by Max Little and colleagues in partnership with the National Centre for Voice and Speech (Denver, CO).

- **Participant Cohort:** 31 individuals (23 clinically diagnosed with Parkinson's, 8 healthy controls).
 - **Total Samples:** 195 sustained phonations (each participant recorded ~6 takes of the vowel sound 'ah/').
 - **Extracted Measurements:** 22 quantitative acoustic features per sample alongside a binary diagnostic label.
 - **Severe Class Imbalance:** 147 samples (75.4%) are Parkinson's cases, and only 48 samples (24.6%) are Healthy controls.

The Imbalance Challenge: A naive algorithm guessing 'Parkinson's' blindly would record an apparent 75.4% accuracy while providing zero screening value. Our pipeline counters this by integrating class reweighting and clinical-grade metrics.

5. Acoustic Vocal Biomarkers Explained in Plain English

Biomarker Category	Key Variables	What It Measures in Human Speech
Fundamental Frequency (Pitch)	MDVP:Fo(Hz), Fhi, Flo	Average, highest, and lowest pitch of the vocal cords. Captures overall vocal range, pitch boundaries, and tonal stability.
Frequency Instability (Jitter)	MDVP:Jitter(%), RAP, PPQ, DDP	Cycle-to-cycle timing variations in pitch. Healthy voices maintain consistent timing; vocal cord micro-tremors cause rapid jitter.
Amplitude Instability (Shimmer)	MDVP:Shimmer, APQ, DDA	Cycle-to-cycle loudness variations. Incomplete vocal cord seal produces rapid, involuntary volume wavering and breathiness.
Noise & Tonal Clarity	HNR, NHR	Harmonics-to-Noise Ratio (HNR) compares pure musical tone to turbulent breathy noise. Lower HNR indicates vocal raspiness.
Nonlinear Signal Dynamics	PPE, RPDE, DFA, Spread1, Spread2	Quantifies turbulence and physical chaos in sound waves. Pitch Period Entropy (PPE) detects irregular pitch drift even when basic pitch averages seem normal.

6. Machine Learning Pipeline & Clinical Evaluation

To ensure that our screening model generalizes reliably to new, unseen patients, all model evaluations were performed using rigorous **Stratified 5-Fold Cross-Validation**. The dataset was partitioned into five subsets that each preserved the 75:25 class distribution. Over five iterations, four subsets trained the classifier while the remaining subset served as a blind test set, rotating until every sample had been evaluated out-of-fold.

Clinical Metrics vs. Raw Accuracy: In a healthcare screening context, overall accuracy is a secondary metric. Medical utility depends fundamentally on two clinical trade-offs:

- **Sensitivity (Recall for Class 1):** The proportion of actual Parkinson's patients correctly identified. A high sensitivity ensures that individuals with neurological impairment are not falsely reassured (minimizing dangerous False Negatives).
- **Specificity (Recall for Class 0):** The proportion of healthy individuals correctly classified as healthy. High specificity prevents unnecessary anxiety, diagnostic follow-ups, and specialist clinical backlogs (minimizing False Positives).
- **Area Under the ROC Curve (AUC):** Measures the classifier's overall discriminative power across all possible decision thresholds, where 0.5 represents a coin flip and 1.0 represents perfect diagnostic separation.

Performance Metric	Baseline Random Forest	Champion XGBoost Model	Clinical Interpretation & Real-World Impact
Average AUC	0.9622	0.9738 (+1.2%)	Exceptional discriminative separation between cohorts across all confidence thresholds.
Sensitivity (Recall PD)	0.9522	0.9524 (Stable)	Catches over 95 out of every 100 individuals showing acoustic signs of Parkinson's.
Specificity (Recall Healthy)	0.6933	0.8356 (+14.2%)	Dramatic reduction in false positive alarms, avoiding unnecessary patient panic.
Derived Accuracy	88.8%	92.3% (+3.5%)	Correctly categorizes ~180 out of the 195 test samples across all folds.

The Engineering Breakthrough (Managing Imbalance): While the initial Random Forest baseline delivered impressive sensitivity (95.2%), its specificity was poor (69.3%), meaning it falsely classified almost one in three healthy participants as having Parkinson's. By advancing to an Extreme Gradient Boosting (XGBoost) architecture with dynamic positive class weighting ($\text{scale_pos_weight} = 48 / 147 = 0.3265$), we heavily penalized misclassifications on the minority healthy cohort. This raised specificity by +14.2 percentage points (to 83.6%) without compromising sensitivity, lifting overall accuracy to 92.3%.

7. Explainable AI & Visual Pitch Deck Portfolio

Opaque 'black-box' algorithms are impractical in healthcare; clinicians must understand *why* a patient was flagged. We generated five publication-ready visual assets (saved in visualizations/) to provide transparent interpretability:

File / Graphic	Visual Format	Stakeholder & Clinical Value
01_shap_summary.png	SHAP Beeswarm Plot	Identifies global feature drivers. Shows that high Pitch Irregularity (PPE) and Spread1 strongly push the model toward a Parkinson's diagnosis.
02_feature_importance.png	Horizontal Bar Chart	Ranks the top 10 acoustic biomarkers by relative mathematical gain, confirming that nonlinear entropy and shimmer drive predictions.
03_case_healthy.png & 03_case_risk.png	Individual Waterfall Plots	Simulates an individual patient diagnostic report. Traces step-by-step why a healthy patient scored low risk, and which biomarkers elevated an at-risk score.
04_metrics_roc.png	Confusion Matrix & ROC Curve	Dual technical verification panel demonstrating low false positive counts and empirical 0.97+ AUC on a 25% holdout test partition.
05_raw_signal_boxplots.png	2x2 Biomarker Boxplots	Overlays individual raw points over cohort distributions for PPE, Jitter, Shimmer, and HNR using a modern blue/coral palette, showing clear separation.

8. Current Limitations & Strategic Roadmap

Key Limitations to Address:

- Cohort Breadth:* While 195 recordings were analyzed, they originated from 31 distinct individuals. Testing on broader multi-speaker populations is essential to prevent subject-level acoustic memorization.
- Acoustic Environment:* Benchmark data was gathered in quiet acoustic chambers. Mobile microphone compression, room reverberation, and ambient noise must be validated.
- Differential Diagnosis:* Other conditions (e.g., laryngitis, essential tremor) can temporarily alter voice and require multi-day baseline tracking.

Future Development Phases:

- Phase 4 (Audio Ingestion & Feature Extraction):** Ingest the Italian and NeuroVoz raw audio repositories (.wav format) using praat-parselmouth, opensmile, and librosa.
- Phase 5 (Mobile Microservice & Edge Deployment):** Export the trained XGBoost model into an embedded Flask API or ONNX runtime for sub-second, on-device mobile inference.

Conclusion & Takeaway: Project NeuroTone demonstrates that vocal acoustic biomarkers alone can differentiate Parkinson's Disease patients from healthy controls with remarkable efficacy (**92.3% accuracy, 0.9738 AUC, 95.2% sensitivity**). By pairing clinical-grade discriminative power with explainable visual interpretations, NeuroTone provides a proven, scalable foundation for accessible, decentralized early neurological screening.